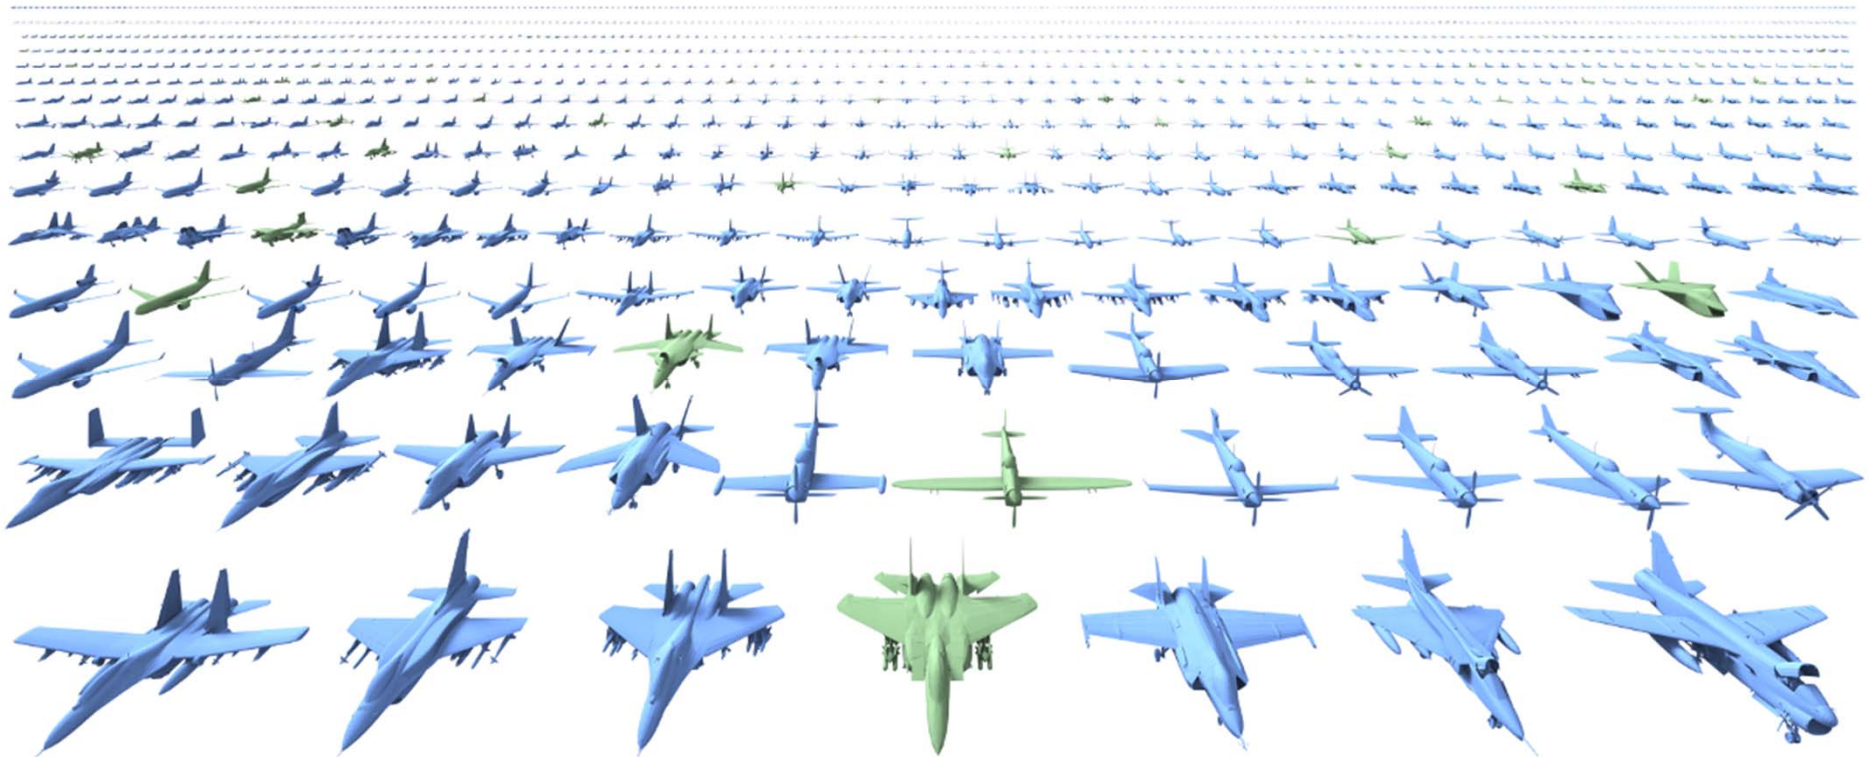


3D Generative models: GANs



Evangelos Kalogerakis –
574/674

How to generate shapes/scenes?

- Encoder-Decoders
 - Case Study: Multi-view decoder
 - Case Study: Implicit decoder
 - Case Study: Patch decoder
 - Case Study: Mesh Decoder
- **Generative Adversarial Networks**
 - Case Study: 3D Volumetric GAN
 - Case Study: Get3D
- Variational Autoencoders
- Autoregressive models
- Diffusion models

GAN basic idea

Learn to generate data (from scratch) based on the underlying distribution of the training set

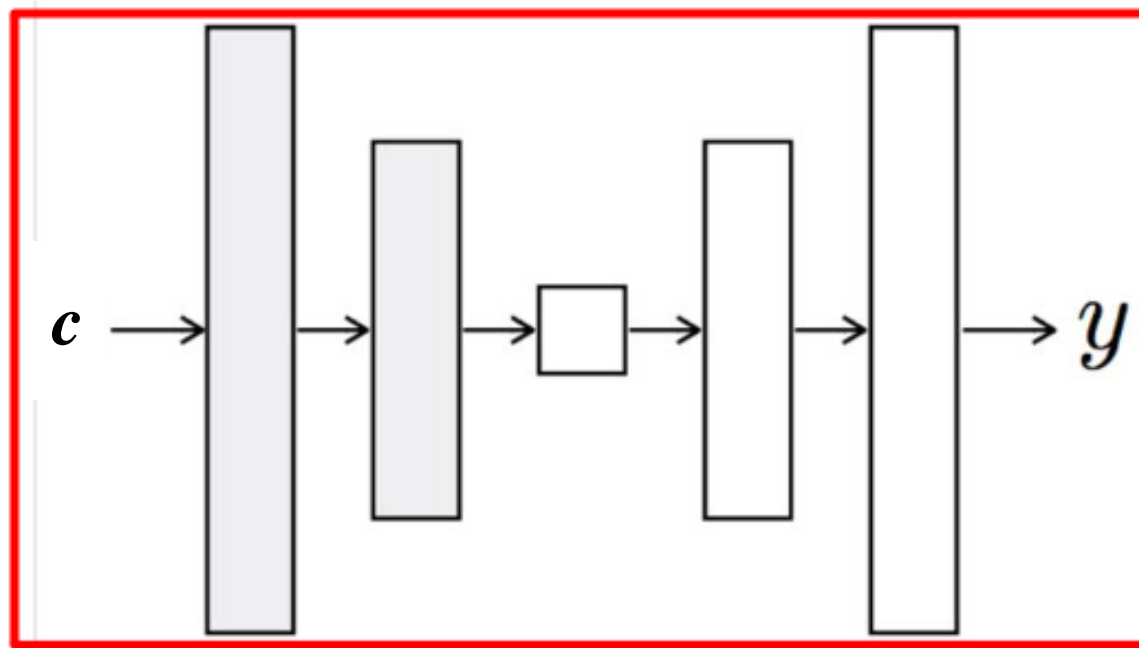
Highly successful in image synthesis

[BigGAN](#) (2018)



GAN basic idea

We need an architecture that can generate data
Previously, we have seen translation networks:

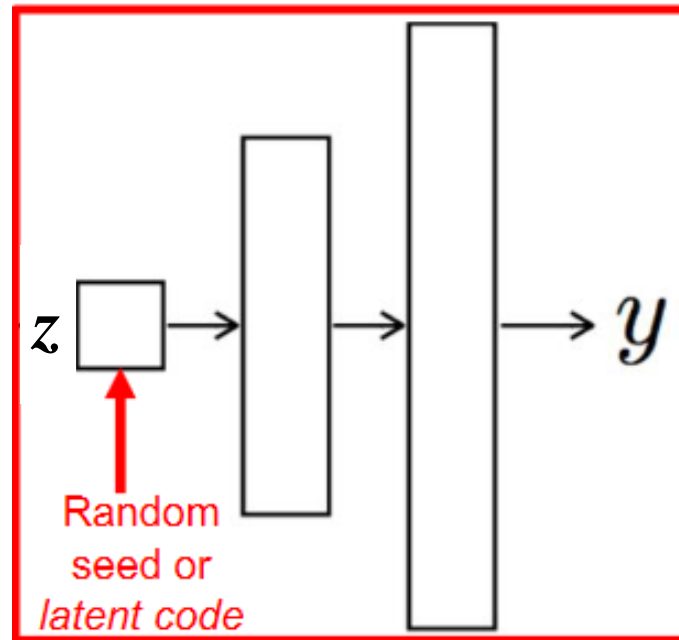


(e.g. input image c , output SDF/mesh/patches y)

GAN basic idea

Start with a random vector (noise) z !

You may use any 3D decoder we discussed before!



GAN basic idea

From: <https://github.com/marian42/shapegan>



Training data from p_{data}



Generated samples p_{model}

We want to learn p_{model} that matches p_{data}

‘Classical’ GAN

Train two networks with **opposing objectives**:

- **Generator:** learns to generate samples
- **Discriminator:** learns to distinguish between generated and real samples

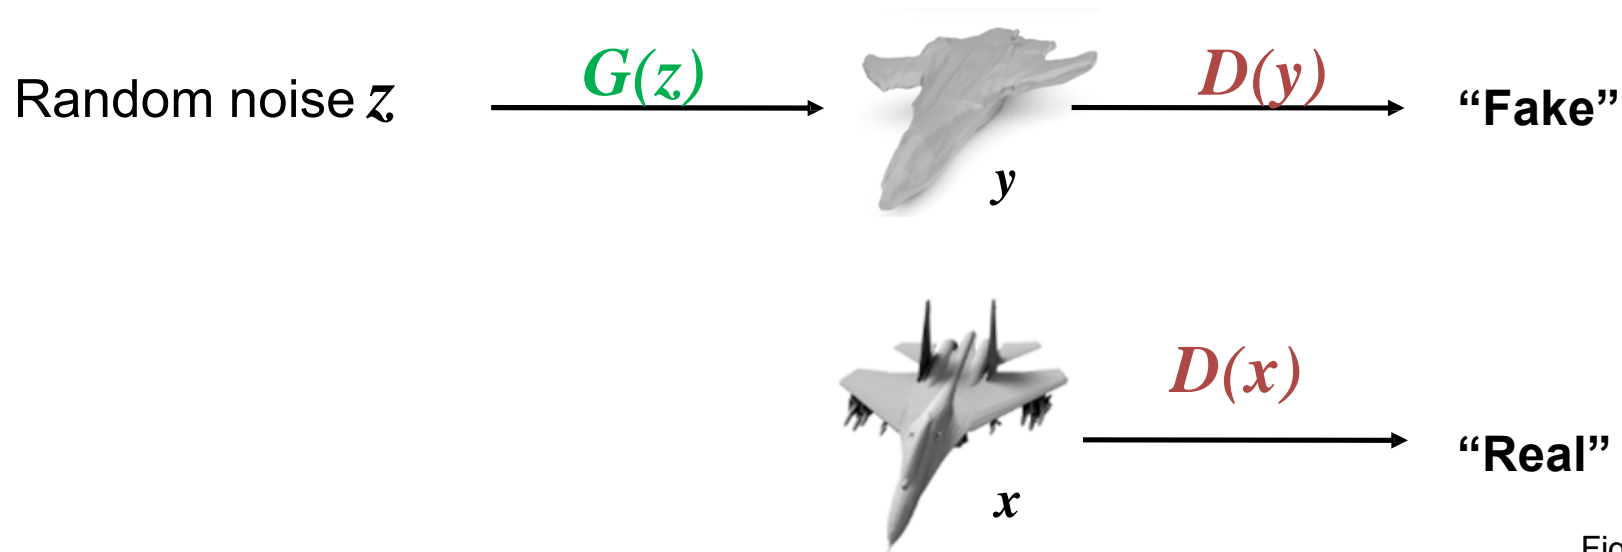


Figure adapted
from [F. Fleuret](#)

I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair,
A. Courville, Y. Bengio, Generative adversarial nets, NIPS 2014

GAN objective

The discriminator D should output the probability that the sample x is real i.e.,

we want $D(x) \approx 1$ for **real data** and

$D(y) = D(G(z)) \approx 0$ for **fake data**

Expected conditional log likelihood for real and generated data:

$$\mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$$

We seed the generator with noise z
drawn from a simple distribution p
(Gaussian or uniform)

GAN objective

$$V(G, D) = \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$$

- The discriminator wants to correctly distinguish real and fake samples:

$$D^* = \arg \max_D V(G, D)$$

- The generator wants to fool the discriminator:

$$G^* = \arg \min_G V(G, D)$$

- Train the generator and discriminator jointly in a *minimax game*

GAN training

$$V(G, D) = \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$$

Alternate between:

- *Gradient ascent* on discriminator:

$$D^* = \arg \max_D V(G, D)$$

- *Gradient descent* on generator (minimize log-probability of discriminator being right):

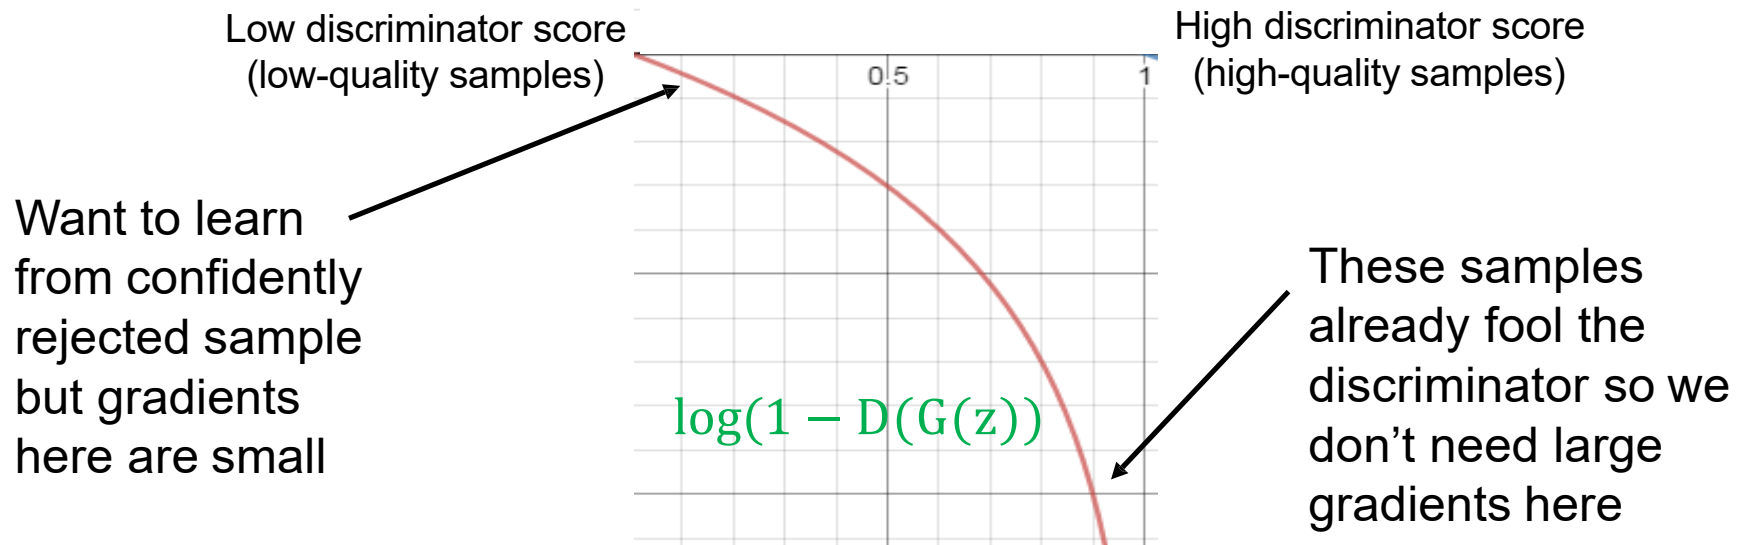
$$\begin{aligned} G^* &= \arg \min_G V(G, D) \\ &= \arg \min_G \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \end{aligned}$$

- In practice, do *gradient ascent* on generator (maximize log-probability of discriminator being wrong):

$$G^* = \arg \max_G \mathbb{E}_{z \sim p} \log(D(G(z)))$$

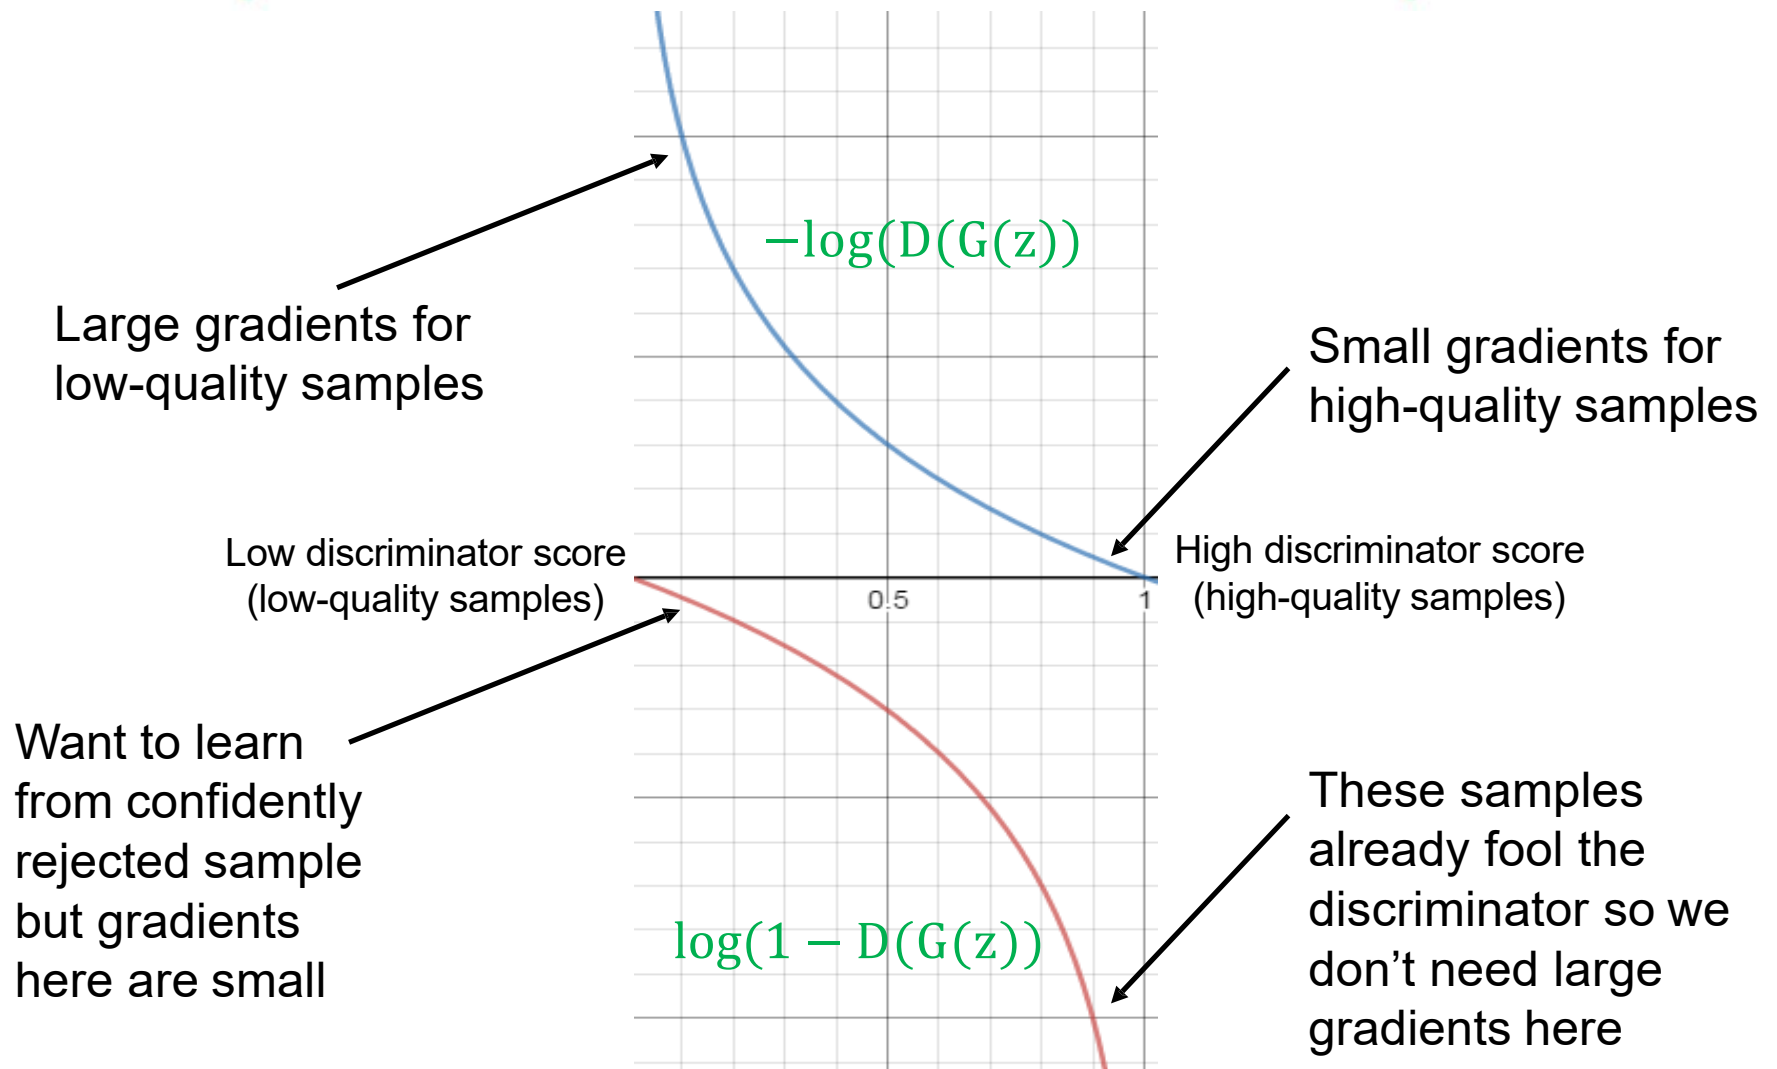
GAN training

$$\min_G \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \text{ vs. } \max_G \mathbb{E}_{z \sim p} \log(D(G(z)))$$



GAN training

$$\min_G \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \text{ vs. } \max_G \mathbb{E}_{z \sim p} \log(D(G(z)))$$



GAN training

Update discriminator

- Repeat the following steps:
 - Sample mini-batch of noise samples $z_1, z_2, z_3, \dots, z_m$ and mini-batch of real samples $x_1, x_2, x_3, \dots, x_m$
 - Update parameters of **D** by stochastic gradient ascent:

$$\frac{1}{m} \sum_m [\log D(x_m) + \log(1 - D(G(z_m)))]$$

Update generator

Sample mini-batch of noise samples $z_1, z_2, z_3, \dots, z_m$

- Update parameters of **G** by stochastic gradient ascent on

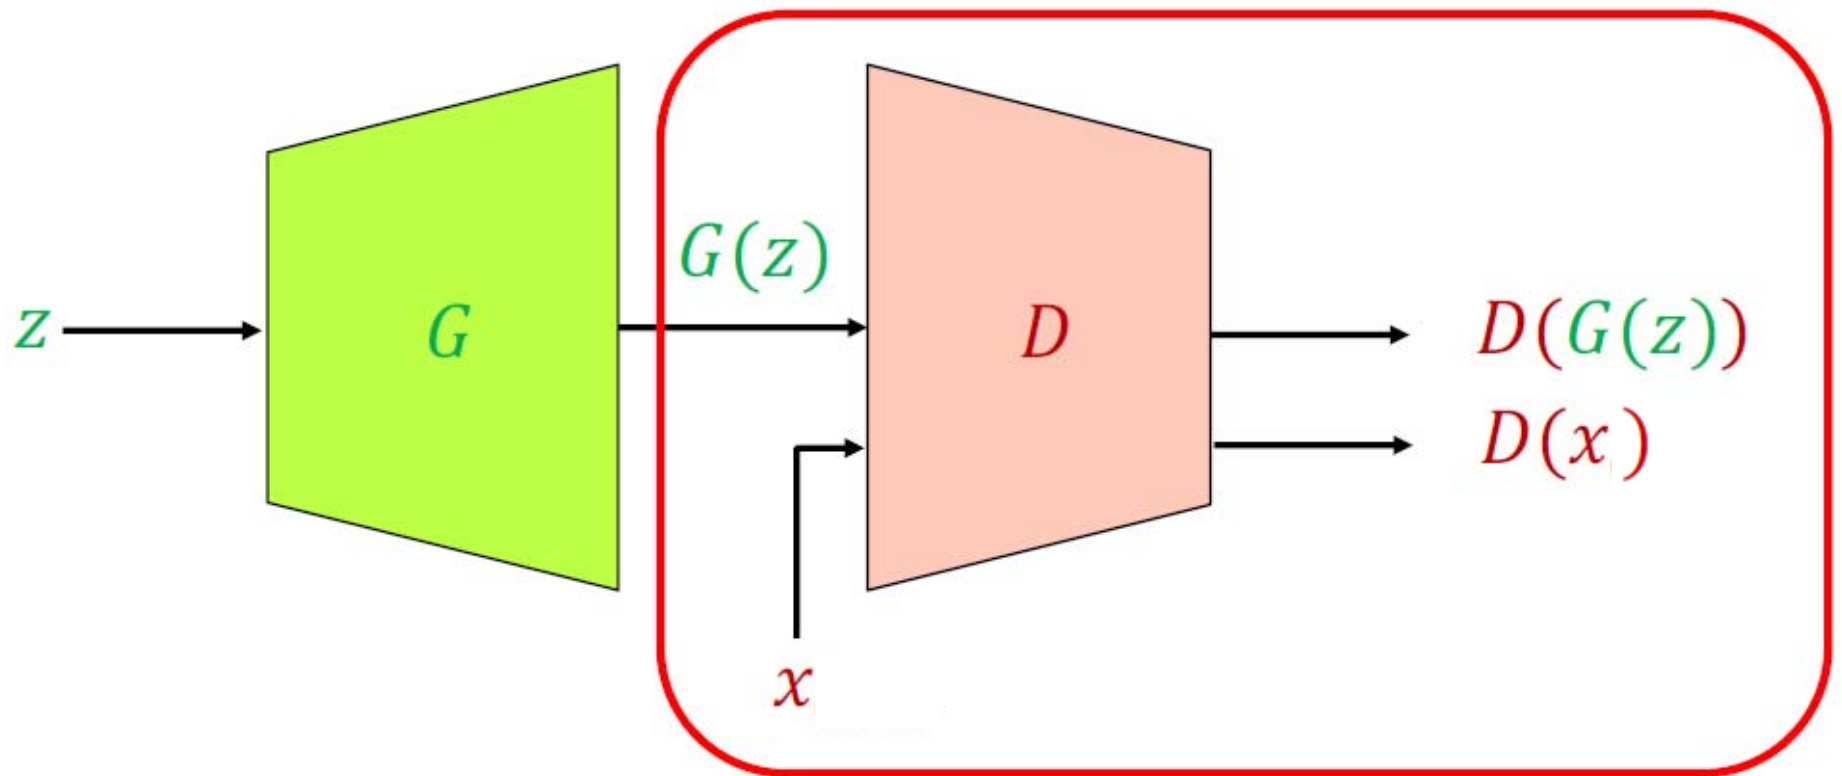
$$\frac{1}{m} \sum_m \log D(G(z_m))$$

... repeat discrim. & gen. training ... until happy with results

GAN conceptual picture

Update discriminator: push $D(x)$ close to 1 and $D(G(z))$ close to 0

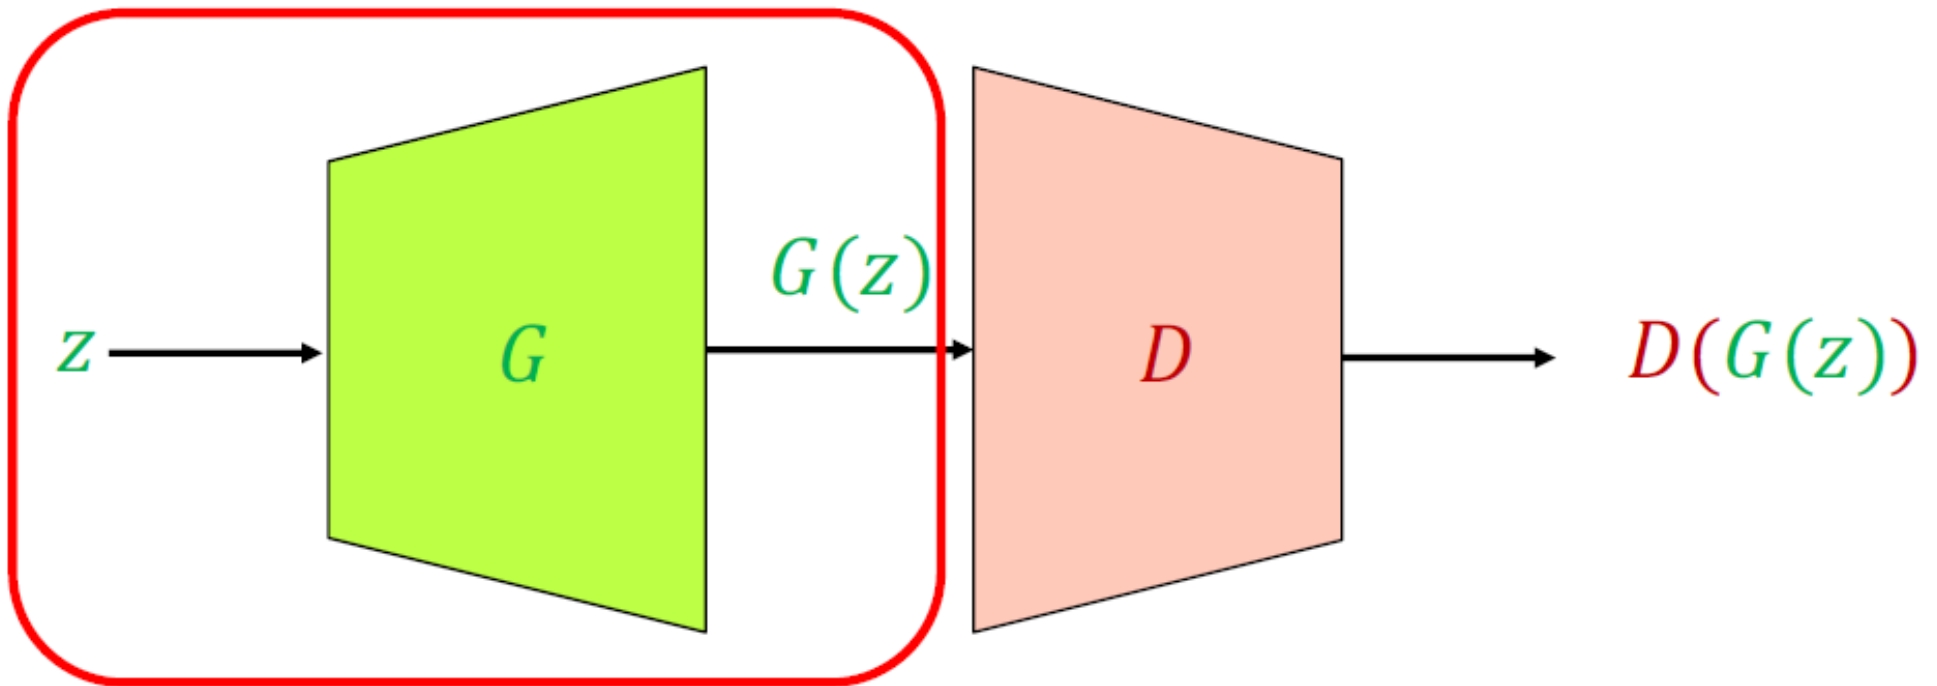
- The generator is a “black box” to the discriminator



GAN conceptual picture

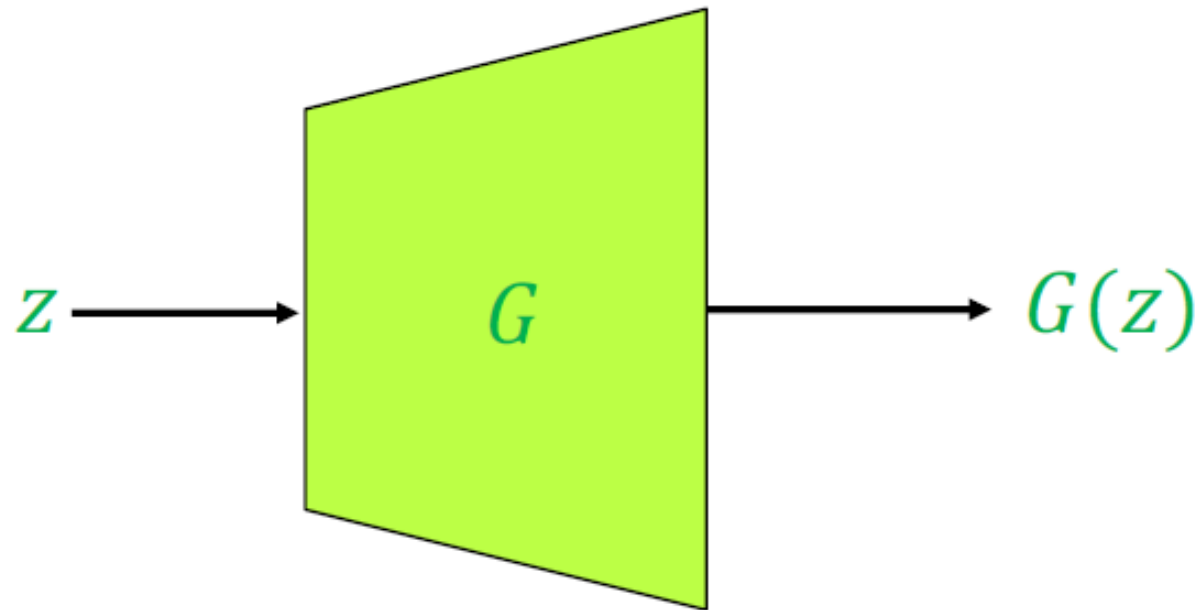
Update generator: increase $D(G(z))$

- Requires back-propagating through the composed generator-discriminator network (i.e., the discriminator cannot be a black box)
- The generator is exposed to real data only via the output of the discriminator (and its gradients)



GAN conceptual picture

Test time:



Problems with GAN training

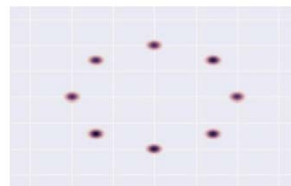
Stability

- Parameters can oscillate or diverge
- Behavior very sensitive to hyper-parameter selection

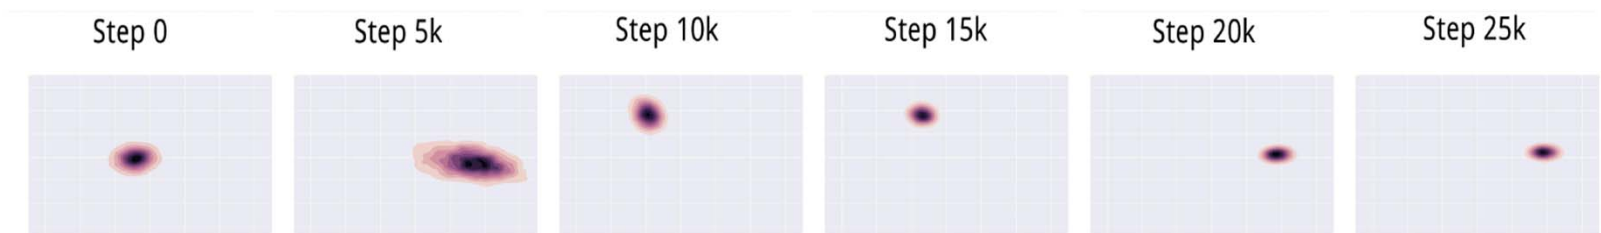
Mode collapse

- Generator ends up modeling only a small subset of the training data.

Target



Output



Wasserstein GAN

Motivated by *Wasserstein or Earth mover's distance*, for comparing distributions

- In practice, simply drop the sigmoid from the discriminator:

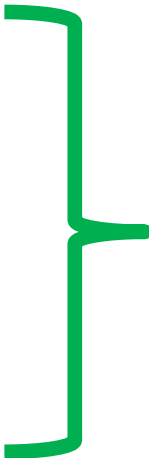
$$\min_G \max_D [\mathbb{E}_{x \sim p_{\text{data}}} D(x) - \mathbb{E}_{z \sim p} D(G(z))]$$

- Need to also clip weights to a fixed range after each gradient update to promote stability

How to generate shapes/scenes?

- Encoder-Decoders

- Case Study: Multi-view decoder
- Case Study: Implicit decoder
- Case Study: Patch decoder
- Case Study: Mesh Decoder



You may use any of these decoders as generators in GANs, together with a corresponding 3D network (multi-view, volumetric, point-based, graph-based) as a discriminator

- **Generative Adversarial Networks**

- Case Study: 3D Volumetric GAN
- Case Study: Get3D

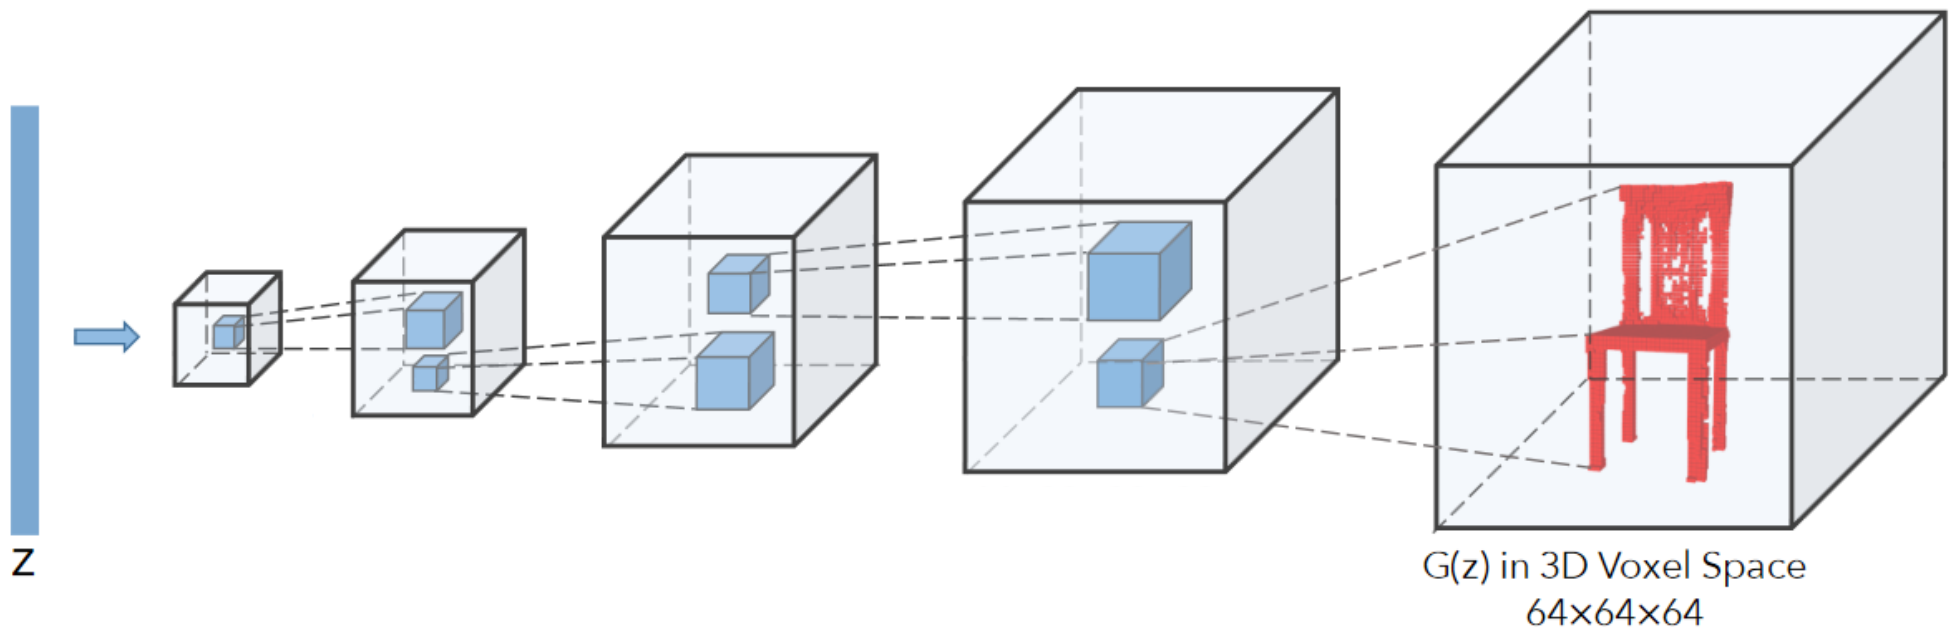
- Variational Autoencoders
- Autoregressive models
- Diffusion models

How to generate shapes/scenes?

- Encoder-Decoders
 - Case Study: Multi-view decoder
 - Case Study: Implicit decoder
 - Case Study: Patch decoder
 - Case Study: Mesh Decoder
- Generative Adversarial Networks
 - **Case Study: 3D Volumetric GAN**
 - Case Study: Get3D
- Variational Autoencoders
- Autoregressive models
- Diffusion models

3D GANs

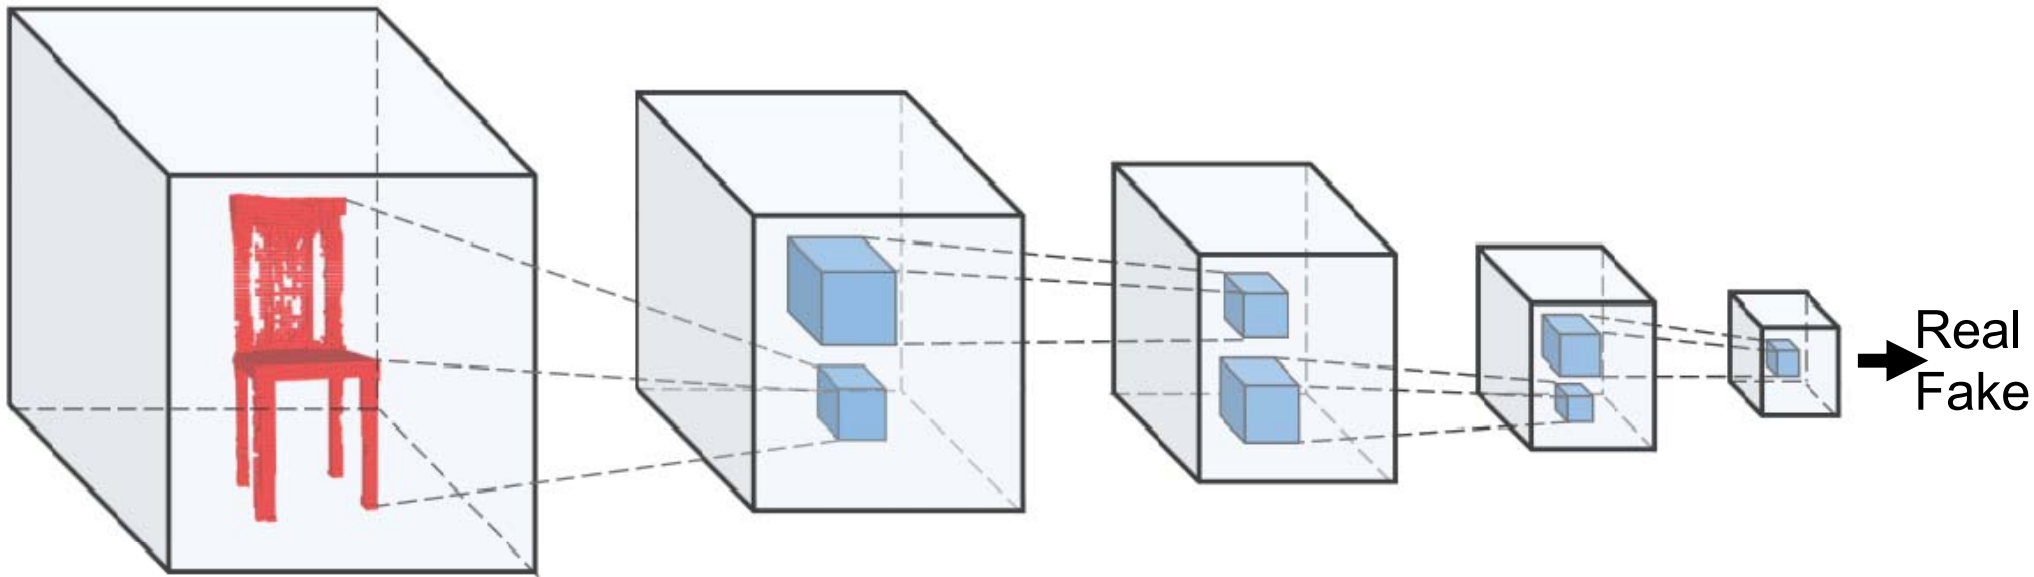
Uses a volumetric decoder to produce a dense grid with a series of transpose 3D convolution layers



Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling, Wu et al. 2016

3D GANs

The discriminator largely mirrors the generator (with a real/fake prediction in the end)



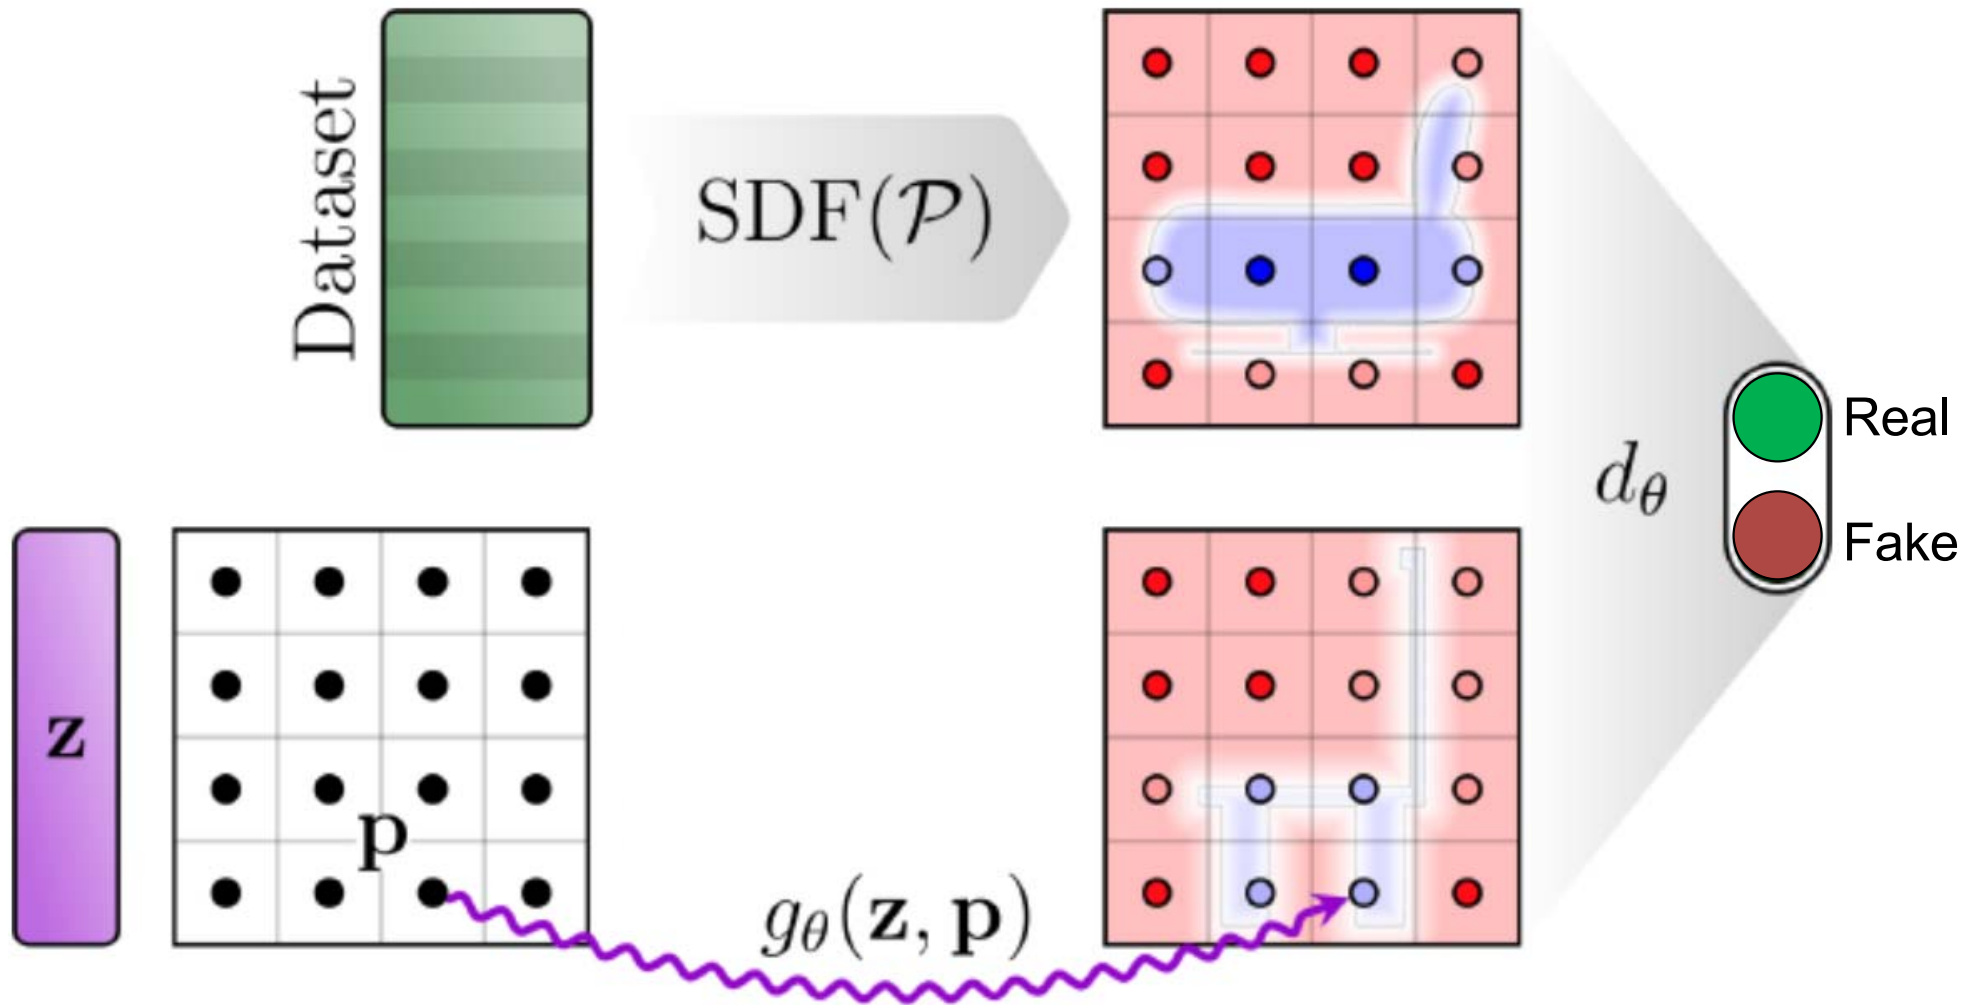
Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling, Wu et al. 2016

Results



Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling, Wu et al. 2016

Extensions



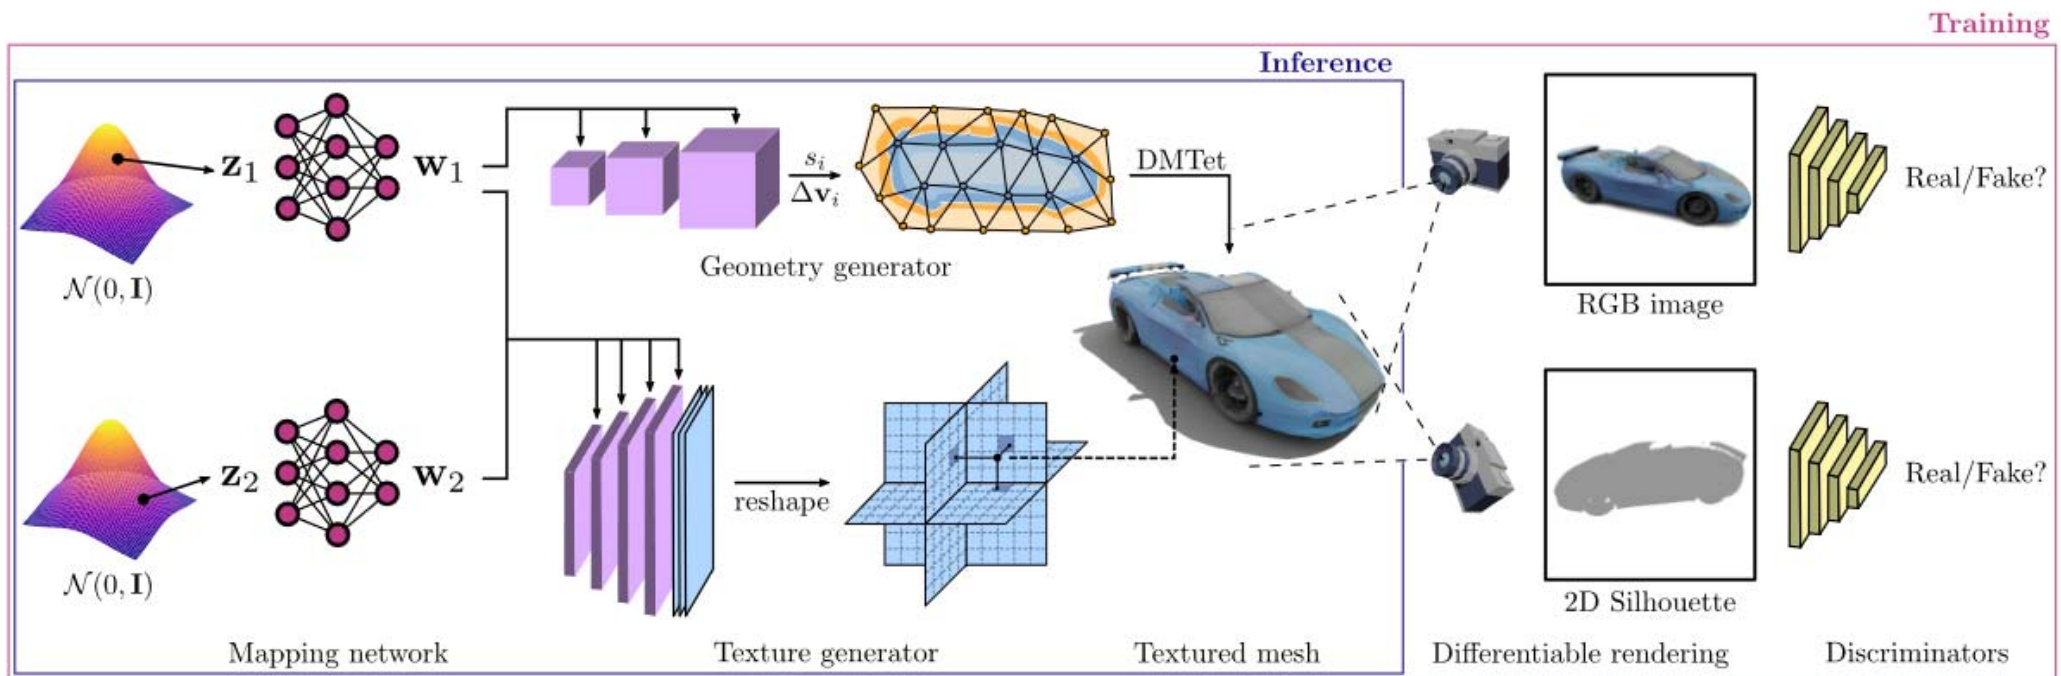
Results



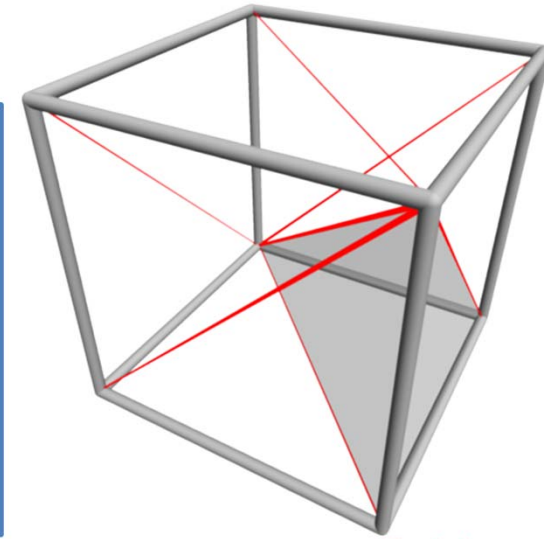
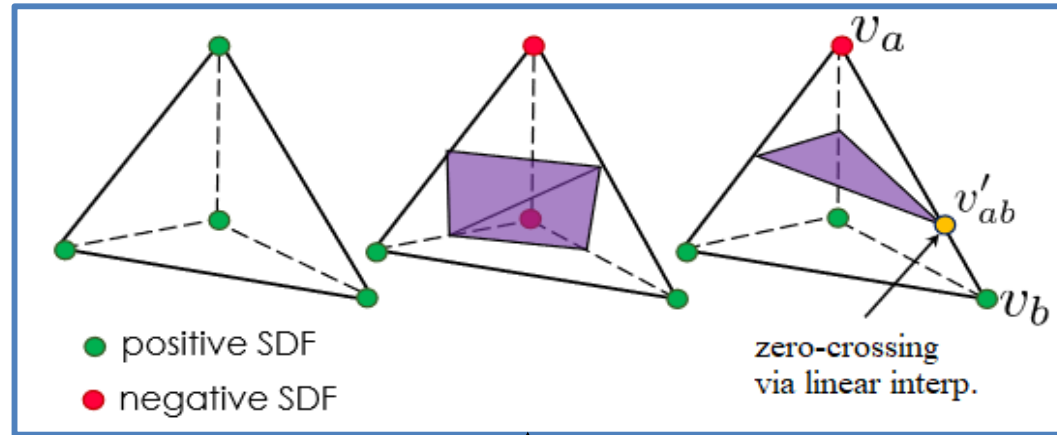
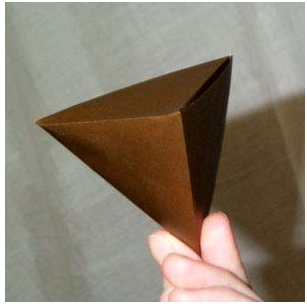
How to generate shapes/scenes?

- Encoder-Decoders
 - Case Study: Multi-view decoder
 - Case Study: Implicit decoder
 - Case Study: Patch decoder
 - Case Study: Mesh Decoder
- Generative Adversarial Networks
 - Case Study: 3D Volumetric GAN
 - **Case Study: Get3D**
- Variational Autoencoders
- Autoregressive models
- Diffusion models

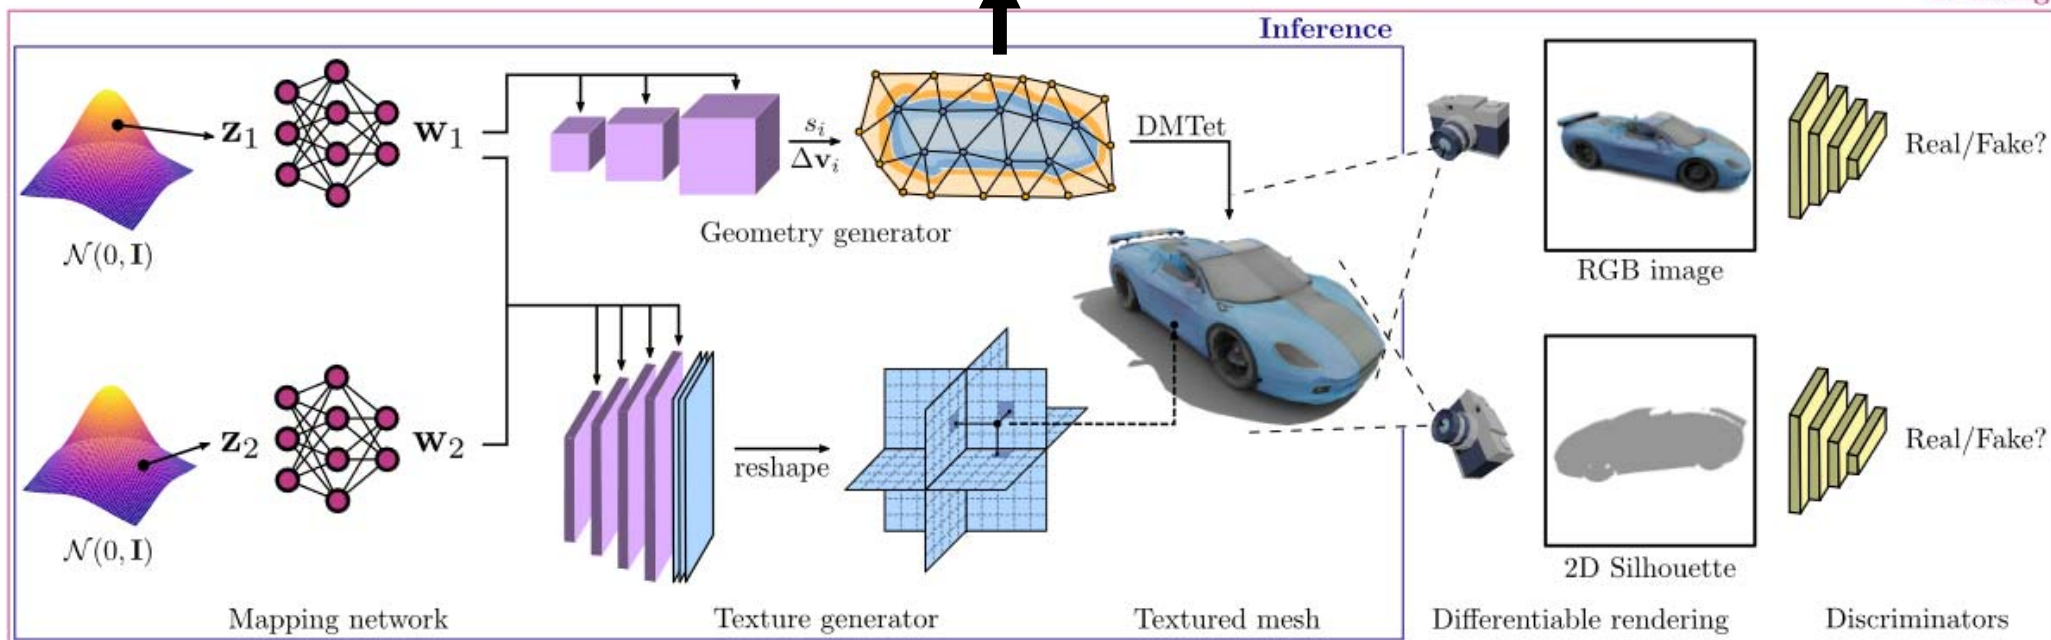
Get3D

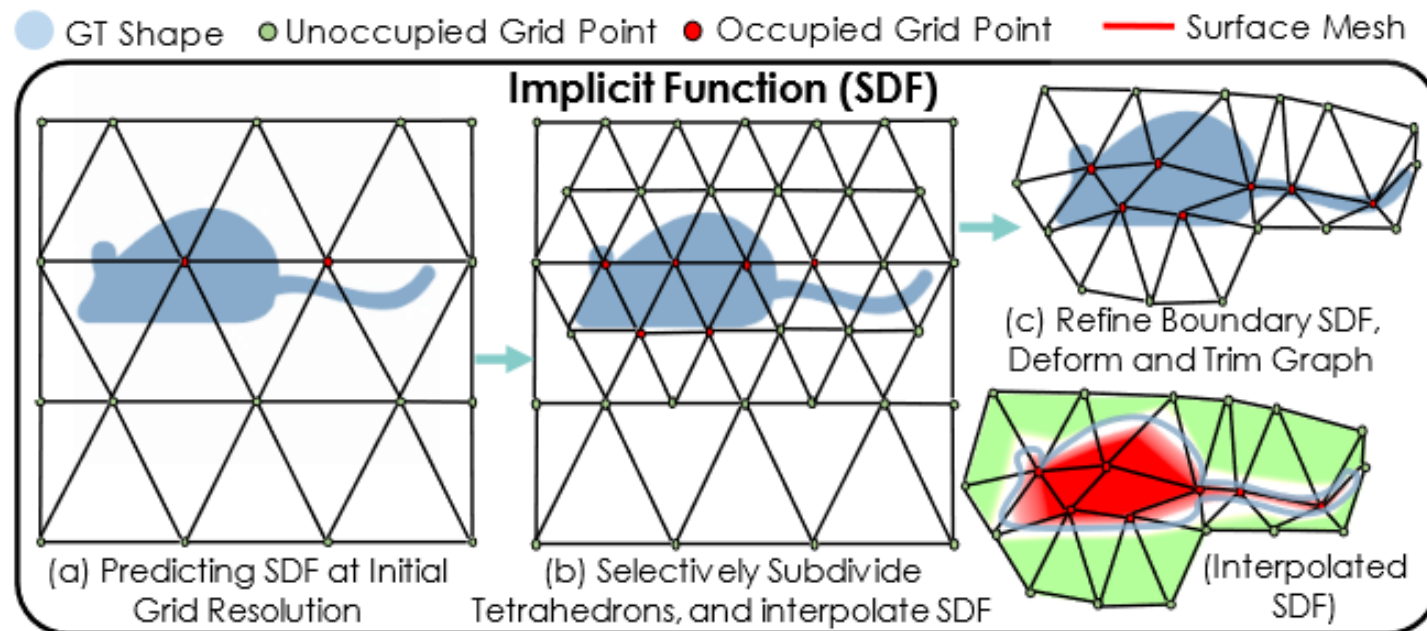


Get3D



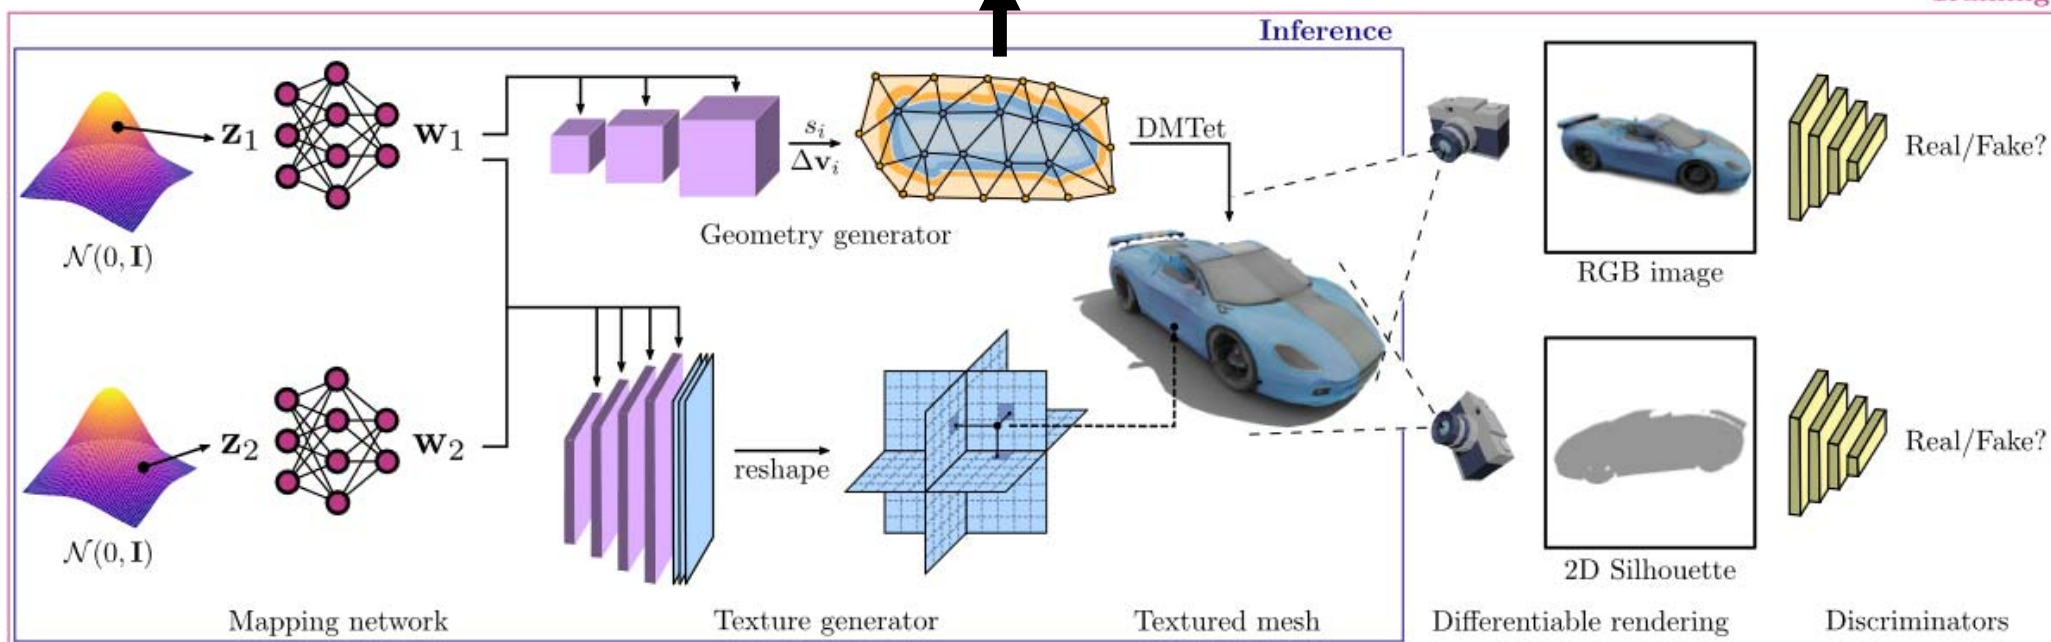
Training





Deep
Marching
Tetrahedra
Shen et al 2021

Training

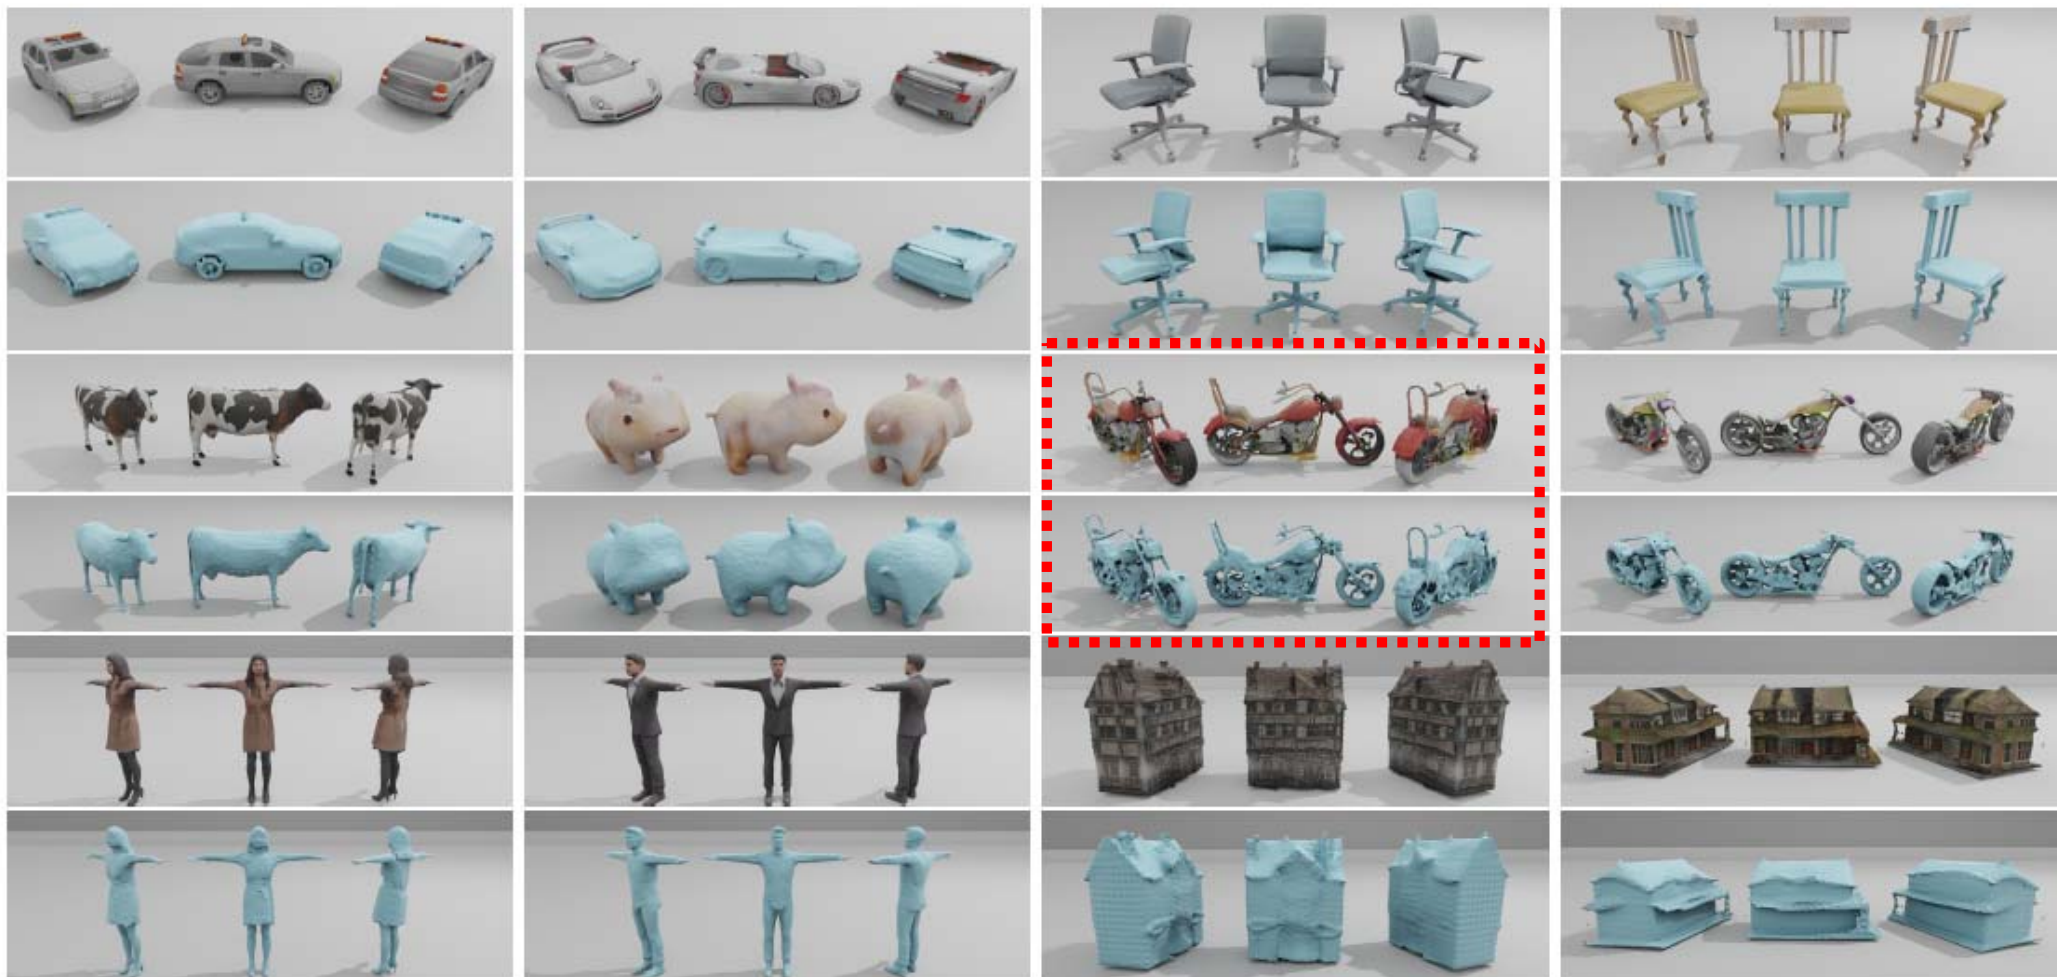


GET3D: A Generative Model of High-Quality 3D Textured Shapes Learned from Images, Gao et al. 2021

Results



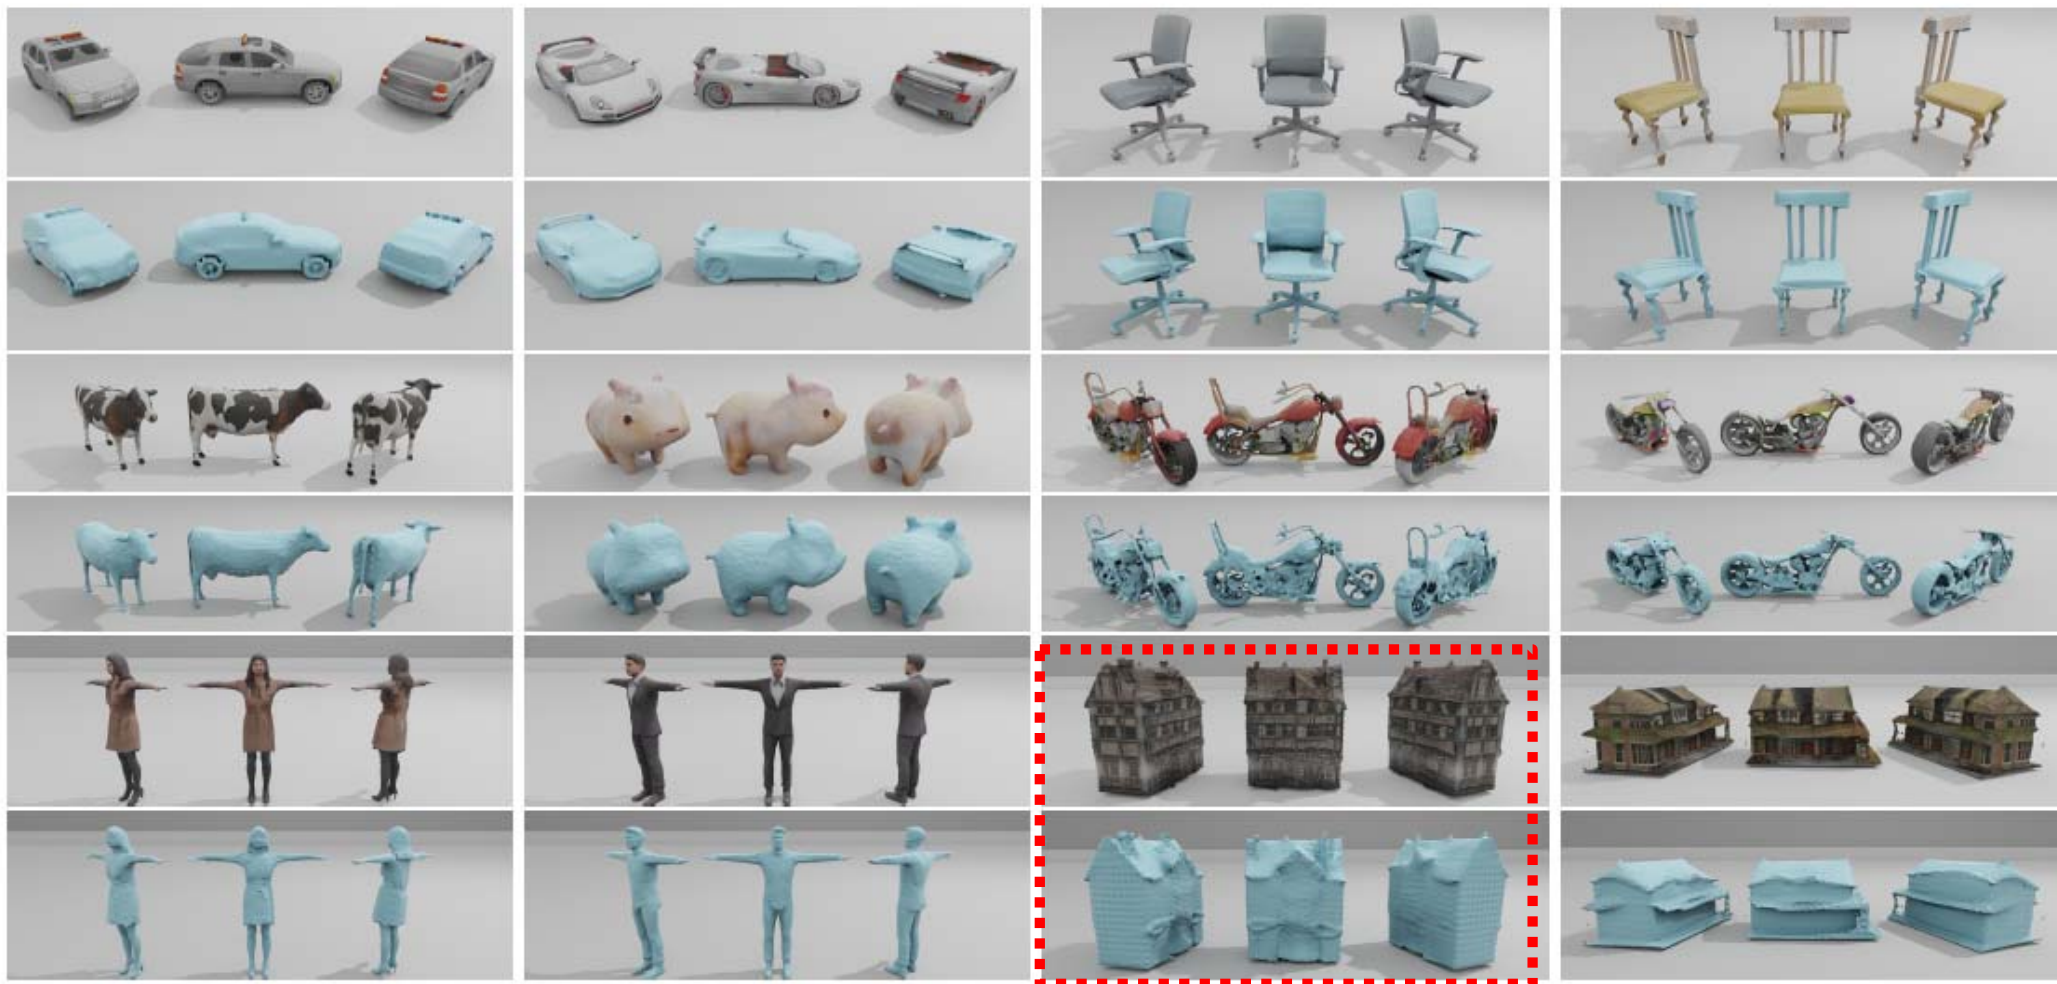
Results

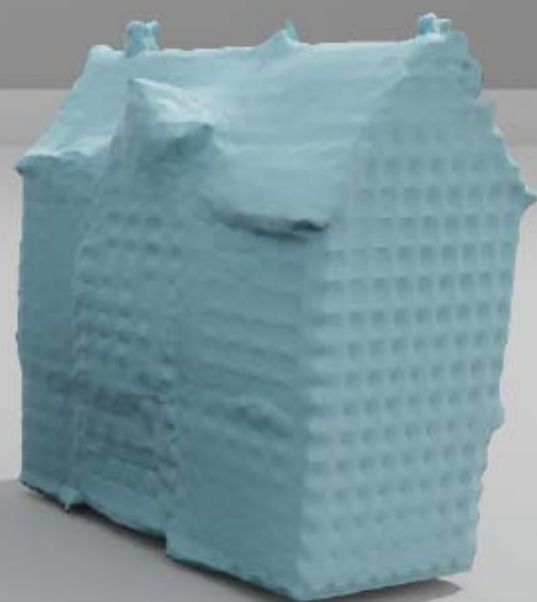
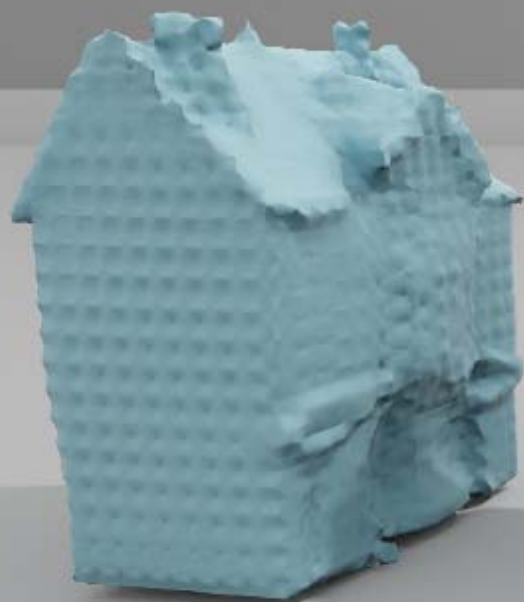


Results

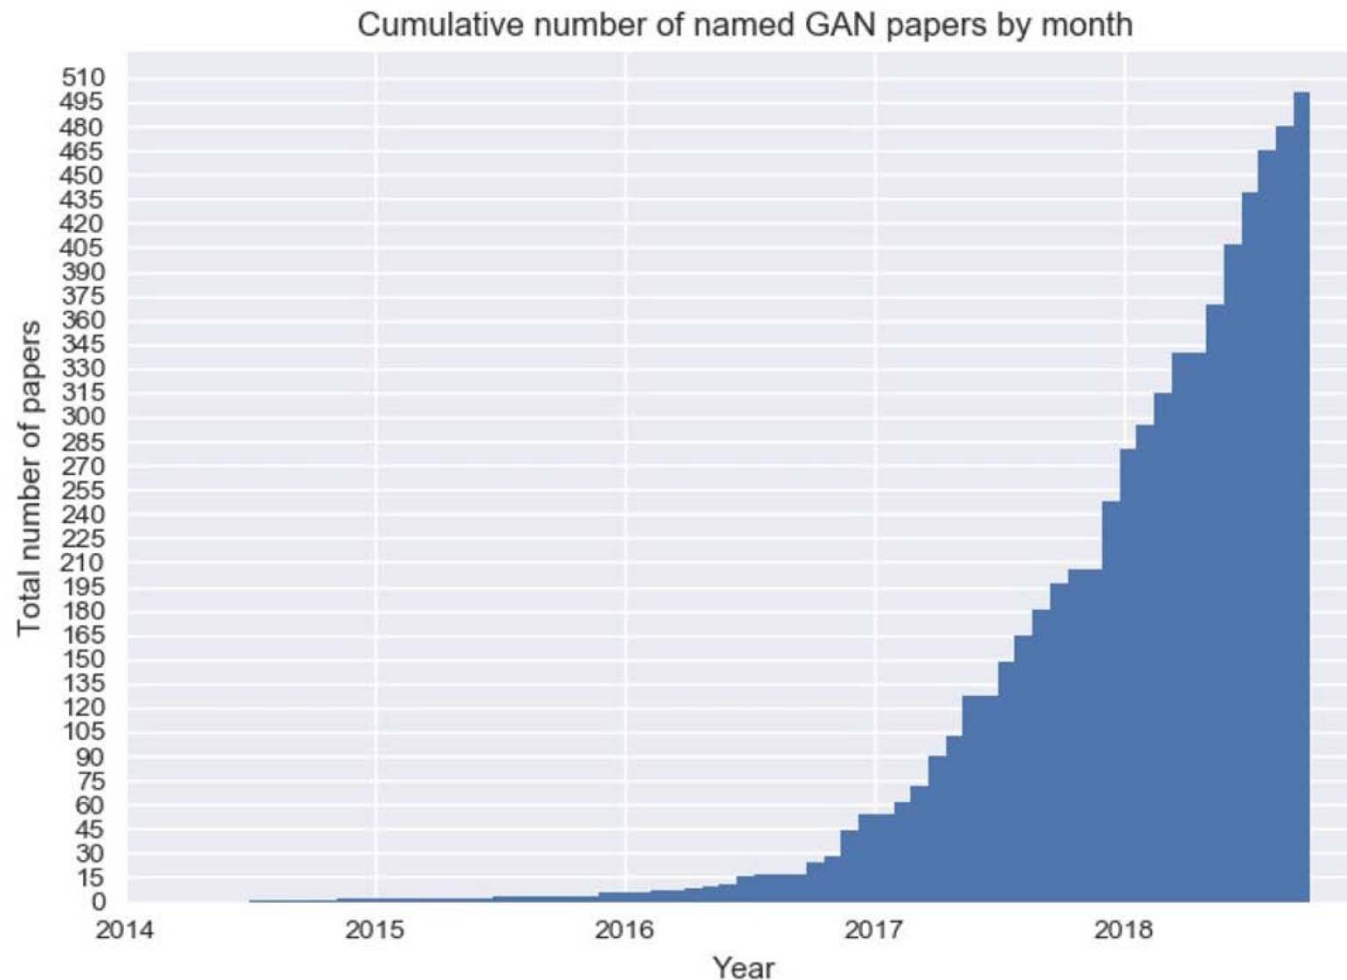


Results





GAN explosion!



<https://github.com/hindupuravinash/the-gan-zoo>

Good GAN papers to read

Resources

- T. Salimans, I. Goodfellow, W. Zaremba, V. Cheung, A. Radford, X. Chen, Improved techniques for training GANs, NIPS 2016
- S. Chintala, E. Denton, M. Arjovsky, M. Mathieu, How to train a GAN? Tips and tricks to make GANs work, 2016
- I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, A. Courville, Improved training of Wasserstein GANs, NIPS 2017
- M. Lucic, K. Kurach, M. Michalski, O. Bousquet, S. Gelly, Are GANs created equal? A large-scale study, NIPS 2018
- P. Isola et al. “Image-to-image translation with conditional adversarial networks.” CVPR 2017.
- J.-Y. Zhu et al. “Unpaired image-to-image translation using cycle-consistent adversarial networks.” ICCV. 2017.
- Zhang, Han, et al. “Self-attention generative adversarial networks”, 2018
- T. Karras, S. Laine, T. Aila, A Style-Based Generator Architecture for Generative Adversarial Networks, CVPR 2019